

Dijkstra's Algorithm

Single source shortest paths
for a directed graph with no
negative edges

Single-Source Shortest Paths

- We want to find the shortest paths between Binghamton and New York City, Boston, and Washington DC. Given a US road map with all the possible routes how can we determine our shortest paths?

Single Source Shortest Paths Problem

- To solve this problem, we may use Floyd's algorithm that finds all pairs shortest paths via dynamic programming. But, this is an overkill, because we have a **single source** now.
- Floyd's algorithm is $\Theta(n^3)$. Can we solve the single source shortest paths problem faster than $\Theta(n^3)$?

Dijkstra's algorithm

- Given a weighted digraph and a vertex s in the graph, find a shortest path from s to an arbitrary node t
- Both for directed and undirected graphs
- No negative edges
- Graph must be connected

Dijkstra's shortest path algorithm

- Dijkstra's algorithm solves the single source shortest path problem in 2 stages.
 - Stage 1: A greedy algorithm computes the shortest *distance* from s to all other nodes in the graph and saves a *data structure*.
 - Stage 2 : Uses the *data structure* to find a shortest path from s to t .

Main idea

- Assume that the shortest distances from the starting node s to the rest of the nodes are
$$d(s, s) \leq d(s, s_1) \leq d(s, s_2) \leq \dots \leq d(s, s_{n-1})$$
- In this case a shortest path from s to s_i may include any of the vertices $\{s_1, s_2, \dots, s_{i-1}\}$ but cannot include any s_j where $j > i$.
- Dijkstra's main idea is to select the nodes and compute the shortest distances in the order $s, s_1, s_2, \dots, s_{n-1}$

Example


$$d(s, s) = 0 \leq$$

$$d(s, s_1) = 5 \leq$$

$$d(s, s_2) = 6 \leq$$

$$d(s, s_3) = 8 \leq$$

$$d(s, s_4) = 15$$



Note: The shortest path from s to s_2 includes s_1 as an intermediate node but cannot include s_3 or s_4 .

Dijkstra's greedy selection rule

- Assume $s_1, s_2 \dots s_{i-1}$ have been selected, and their shortest distances have been stored in **Solution**
- Select node s_i and save $d(s, s_i)$ if s_i has the shortest distance from s on a path that may include only $s_1, s_2 \dots s_{i-1}$ as intermediate nodes. We call such paths *special*
- To apply this selection rule efficiently, we need to maintain for each unselected node v the *distance of the shortest special path* from s to v , $D[v]$.

Application Example

$Solution = \{(s, 0)\}$

$D[s_1]=5$ for path $[s, s_1]$

$D[s_2]=\infty$ for path $[s, s_2]$

$D[s_3]=10$ for path $[s, s_3]$

$D[s_4]=15$ for path $[s, s_4]$.

$Solution = \{(s, 0), (s_1, 5)\}$

$D[s_2]=6$ for path $[s, s_1, s_2]$

$D[s_3]=9$ for path $[s, s_1, s_3]$


$D[s_4]=15$ for path $[s, s_4]$

$Solution = \{(s, 0), (s_1, 5), (s_2, 6)\}$

$D[s_3]=8$ for path $[s, s_1, s_2, s_3]$

$D[s_4]=15$ for path $[s, s_4]$

$Solution = \{(s, 0), (s_1, 5), (s_2, 6), (s_3, 8), (s_4, 15)\}$



Implementing the selection rule

- Node *near* is selected and added to *Solution* if $D(near) \leq D(v)$ for any $v \notin \text{Solution}$.

$Solution = \{(s, 0)\}$


$D[s_1]=5 \leq D[s_2]=\infty$

$D[s_1]=5 \leq D[s_3]=10$

$D[s_1]=5 \leq D[s_4]=15$


Node s_1 is selected

$Solution = \{(s, 0), (s_1, 5)\}$



Updating $D[]$

- After adding *near* to *Solution*, $D[v]$ of all nodes $v \notin \text{Solution}$ are updated if there is a shorter special path from s to v that contains node *near*, i.e., if $(D[\text{near}] + w(\text{near}, v) < D[v])$ then $D[v] = D[\text{near}] + w(\text{near}, v)$



Example: Updating D

Solution = $\{(s, 0)\}$

$D[s_1]=5, D[s_2]=\infty, D[s_3]=10, D[s_4]=15.$

Solution = $\{(s, 0), (s_1, 5)\}$

$D[s_2]=D[s_1]+w(s_1, s_2)=5+1=6,$

$D[s_3]=D[s_1]+w(s_1, s_3)=5+4=9,$


$D[s_4]=15$

Solution = $\{(s, 0), (s_1, 5), (s_2, 6)\}$

$D[s_3]=D[s_2]+w(s_2, s_3)=6+2=8,$

$D[s_4]=15$

Solution = $\{(s, 0), (s_1, 5), (s_2, 6), (s_3, 8), (s_4, 15)\}$



Dijkstra's Algorithm for Finding the Shortest Distance from a Single Source

Dijkstra(G, s)

1. **for each** $v \in V$
2. **do** $D[v] \leftarrow \infty$
3. $D[s] \leftarrow 0$
4. $MH \leftarrow \text{make-MH}(D, V)$ // MH : MinHeap
5. **while** $MH \neq \emptyset$
6. $near \leftarrow MH.\text{extractMin}()$
7. **for each** $v \in \text{Adj}(near)$
8. **if** $D[v] > \mathbf{D[near]} + w(near, v)$
9. **then** $D[v] \leftarrow \mathbf{D[near]} + w(near, v)$
10. $MH.\text{decreaseDistance}(D[v], v)$
11. **return** the label $D[u]$ of each vertex u

Time Complexity Analysis

```
1. for each  $v \in V$ 
2.   do  $D[v] \leftarrow \infty$ 
3.  $D[s] \leftarrow 0$ 
4.  $MH \leftarrow \text{make-MH}(D, V)$ 

5. while  $MH \neq \emptyset$ 
6.   do  $near \leftarrow MH.\text{extractMin}()$ 
7.     for each  $v \in \text{Adj}(near)$ 
8.       if  $D[v] > D[near] + w(near, v)$ 
9.         then  $D[v] \leftarrow D[near] + w(near, v)$ 
10.       $MH.\text{decreaseDistance}(D[v], v)$ 
11. return the label  $D[u]$  of each vertex  $u$ 
```

Assume a node in MH can be accessed in $O(1)$

Using Heap implementation

Lines 1 - 4 run in $O(V)$

Max Size of MH is $|V|$

(5) Loop = $O(V)$

(6) $O(\lg V)$

(5+6) $O(V \lg V)$

(7, 8, 9) are $O(1)$ and executed $O(E)$ times in total

(10) Decrease- Key operation on the heap takes $O(\lg V)$ time, and is executed $O(E)$ times in total

→ $O(E \lg V)$

So total time is $O(V \lg V + E \lg V)$
= $O(E \lg V)$

Alternative way to implement Dijkstra's algorithm

- Use an array instead of a MinHeap
- Time Complexity
 - $O(V)$ to extract min
 - $O(1)$ for decreaseDistance
 - Thus, $O(V^2)$ in total